

1


```
corn_data = read.table("data/cornnit.txt", header=T)
```

```
head(corn_data)
```

	yield	nitrogen
1	115	0
2	128	75
3	136	150
4	135	300
5	97	0
6	150	75

(i) (1 pt) What does the scatter plot of yields versus nitrogen suggest?

```
plot(y=corn_data$yield, x=corn_data$nitrogen, xlab="Nitrogen Application", ylab="Corn Yields")
```



The scatter plot suggests a non linear relationship between nitrogen application and corn yields because the scatter plot does not form a straight line and is instead high slope at the start then starts to level off towards the end.

(ii) (1 pt) Fit separate linear first-order and second-order models with yield as the response and nitrogen as the predictor.

```
fit1 = lm(yield ~ nitrogen, data=corn_data)
summary(fit1)
```

Call:

```
lm(formula = yield ~ nitrogen, data = corn_data)
```

Residuals:

Min	1Q	Median	3Q	Max
-60.439	-10.939	1.534	14.082	29.697

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	107.43864	4.66622	23.02	< 2e-16 ***
nitrogen	0.17730	0.03377	5.25	4.71e-06 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 20.53 on 42 degrees of freedom

Multiple R-squared: 0.3962, Adjusted R-squared: 0.3818

F-statistic: 27.56 on 1 and 42 DF, p-value: 4.713e-06

```
fit2 = lm(yield ~ nitrogen + I(nitrogen^2), data=corn_data)
summary(fit2)
```

Call:

```
lm(formula = yield ~ nitrogen + I(nitrogen^2), data = corn_data)
```

Residuals:

Min	1Q	Median	3Q	Max
-47.162	-7.092	0.368	10.058	26.571

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	94.1620969	4.4067757	21.368	< 2e-16 ***
nitrogen	0.5794901	0.0800285	7.241	7.54e-09 ***
I(nitrogen^2)	-0.0014831	0.0002787	-5.321	3.97e-06 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


Residual standard error: 15.98 on 41 degrees of freedom

Multiple R-squared: 0.6428, Adjusted R-squared: 0.6254

F-statistic: 36.9 on 2 and 41 DF, p-value: 6.817e-10

(iii) (2 pt) Is the second-order model a better fit?

```
par(mfrow=c(1,2))
plot(fit1, 1, main="1st-order linear")
plot(fit2, 1, main="2nd-order linear")
```



```
anova(fit1, fit2, test="F")
```

Analysis of Variance Table

Model 1: yield ~ nitrogen

Model 2: yield ~ nitrogen + I(nitrogen^2)

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	42	17699				
2	41	10470	1	7229.6	28.312	3.974e-06 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

According to the anova test there is a significant relationship between yield and the extra predictor nitrogen^2 . Thus, the second-order model is a better fit than the first-order model.

(iv) (2 pt) What can you conclude about the relationship between yield and nitrogen?

This the relationship between yield and nitrogen is non linear because the second-order model is a better fit since the nitrogen^2 predictor had a statistically significant relationship with the response variable yield.

2

```
diabetes = read.csv("data/diabetes.csv")

head(diabetes)
```

	age	ratio	gender	frame	diabetic
1	46	3.6	female	medium	0
2	29	6.9	female	large	1
3	58	6.2	female	large	0
4	67	6.5	male	large	0
5	64	8.9	male	medium	0
6	34	3.6	male	large	0

(i) (1 pt) Fit a multiple linear regression with ratio as the response and age, gender, and frame as predictors.

```
mult = lm(ratio ~ age + gender + frame, data = diabetes)
summary(mult)
```

Call:

```
lm(formula = ratio ~ age + gender + frame, data = diabetes)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-2.9760	-1.1535	-0.2371	0.8712	14.6783

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.267848	0.347959	12.265	< 2e-16 ***

```

age          0.011385    0.005474    2.080    0.0382 *
gendermale   0.207224    0.175980    1.178    0.2397
framemedium -0.226777    0.213246   -1.063    0.2882
framesmall  -0.970578    0.244779   -3.965    8.75e-05 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Residual standard error: 1.674 on 385 degrees of freedom
Multiple R-squared: 0.07583, Adjusted R-squared: 0.06623
F-statistic: 7.897 on 4 and 385 DF, p-value: 3.985e-06

(ii) (2 pt) Is there a significant difference in the ratio of HDL between males and females?

There is not a significant difference in the ratio of HDL between males and females because the male coefficient has a p-value of 0.2397 which is greater than $\alpha = 0.1$ and additionally $0 \in [0.207224 - (1.64)0.17598, 0.207224 + (1.64)0.17598]$. This means that the gender male coefficient does not differentiate significantly from baseline (gender = female)

(iii) (2pt) Report and interpret the coefficients for frame.

```
summary(mult)$coefficients
```


	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.26784826	0.34795934	12.265365	2.024249e-29
age	0.01138544	0.00547431	2.079795	3.820548e-02
gendermale	0.20722412	0.17597989	1.177544	2.397057e-01
framemedium	-0.22677652	0.21324632	-1.063449	2.882451e-01
framesmall	-0.97057840	0.24477916	-3.965119	8.746412e-05

The framesmedium predictor has a coefficients of -0.22677652 with a std of 0.21324632 and p-value of 0.2882451. This means that the difference between a medium frame and a large frame is not statistically significant because the p-value is greater than $\alpha = 0.1$ and $0 \in [-0.22677652 - (1.64)0.21324632, -0.22677652 + (1.64)0.21324632]$.

The framesmall predictor has a coefficients of -0.97057840 with a std of 0.24477916 and p-value of $8.746412e - 05$. This means that the difference between a small frame and a large frame is statistically significant because the p-value is smaller than $\alpha = 0.1$ and $0 \notin [-0.97057840 - (1.64)0.24477916, -0.97057840 + (1.64)0.24477916]$. More specifically, the average change in the response mean from a change from large frame to small frame is -0.97057840 .

(iv) (1 pt) Is there an interaction between gender and frame based on the interaction plot?

```
interaction.plot(x.factor = diabetes$gender, trace.factor = diabetes$frame, response=diabetes$ratio)
```



There is an interaction between gender and frame based on the interaction plot because the lines are not parallel. More specifically, it appears there may be an interaction between frame and gender because large/medium frame increases the HDL ratio when going from female to male but the small frame decreases the HDL ratio when going from female to male.

(v) (1 pt) If yes, is the interaction effect between gender and frame significant at $\alpha = 0.10$?

```
int_model = lm(ratio ~ age + gender * frame + gender + frame, data = diabetes)
summary(int_model)
```

Call:

```
lm(formula = ratio ~ age + gender * frame + gender + frame, data = diabetes)
```

Residuals:

Min	1Q	Median	3Q	Max
-3.0129	-1.1097	-0.2470	0.8891	14.6982

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.081987	0.383829	10.635	<2e-16 ***
age	0.011213	0.005471	2.050	0.0411 *
gendermale	0.536402	0.335492	1.599	0.1107
framemedium	-0.052085	0.303152	-0.172	0.8637
framesmall	-0.626892	0.331606	-1.890	0.0594 .
gendermale:framemedium	-0.277684	0.421682	-0.659	0.5106
gendermale:framesmall	-0.785332	0.485208	-1.619	0.1064

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.673 on 383 degrees of freedom

Multiple R-squared: 0.08229, Adjusted R-squared: 0.06791

F-statistic: 5.724 on 6 and 383 DF, p-value: 1.021e-05

```
anova(mult, int_model)
```

Analysis of Variance Table

Model 1: ratio ~ age + gender + frame

Model 2: ratio ~ age + gender * frame + gender + frame

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	385	1079.1				
2	383	1071.6	2	7.5455	1.3484	0.2609

The interaction affect appears not to be significant at $\alpha = 0.10$ because the anova table resulted in a p-value of 0.2609 which is greater than our α . This means that the extra predictors do not have statistically significant relationship with the response mean HDL ratio.

3


```
job_data = read.table("data/jobprof.txt", header=T)
```

```
head(job_data)
```

	test1	test2	test3	test4	JBS
1	88	86	110	100	87
2	80	62	97	99	100
3	96	110	107	103	103
4	76	101	117	93	95
5	80	100	101	95	88
6	73	78	85	95	84

(i) (1 pt) Calculate the correlation matrix of the test scores.

```
X = subset(job_data, select=-c(JBS))  
corrplot.mixed(cor(X), order="AOE")
```



(ii) (1 pt) Does the correlation matrix suggest a possible multicollinearity problem for predicting job proficiency score (JPS) from the test scores?

The correlation matrix suggests possible multicollinearity problems because the correlation coefficient between test1 and test4 is 0.9 which is close to 1.

(iii) (1 pt) Fit a multiple linear regression model with JPS as the response and the test scores as the predictors.

```
model = lm(JBS ~ test1 + test2 + test3 + test4, data=job_data)  
summary(model)
```


Call:

```
lm(formula = JBS ~ test1 + test2 + test3 + test4, data = job_data)
```

Residuals:

Min	1Q	Median	3Q	Max
-8.7283	-2.6769	-0.7255	3.1096	9.3900

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	100.88142	30.03086	3.359	0.00312 **
test1	0.84060	0.21337	3.940	0.00081 ***
test2	-0.19182	0.09142	-2.098	0.04879 *
test3	-0.04574	0.07258	-0.630	0.53570
test4	-0.58529	0.40537	-1.444	0.16427

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.212 on 20 degrees of freedom

Multiple R-squared: 0.8014, Adjusted R-squared: 0.7617

F-statistic: 20.17 on 4 and 20 DF, p-value: 8.612e-07

(iv) (1 pt) Compute the variance inflation factors for each of the model parameters in (iii).

```
vif(model)
```

test1	test2	test3	test4
15.172928	3.042143	1.391685	11.384328

(v) (1 pt) What does the VIF suggest?

test1 and test4 have VIF scores above 10 which suggests there is multicollinearity within the data.

(vi) (2 pt) Use a stepwise regression with AIC criterion to select the best model for predicting JPS.

```
library(MASS) ## this package is required for the stepAIC function
both_sel <- stepAIC(model, direction = "both")
```

Start: AIC=86.97

JBS ~ test1 + test2 + test3 + test4

	Df	Sum of Sq	RSS	AIC
- test3	1	10.79	554.10	85.462

<none>			543.31	86.970	
- test4	1	56.63	599.94	87.449	
- test2	1	119.59	662.90	89.944	
- test1	1	421.65	964.96	99.330	

Step: AIC=85.46

JBS ~ test1 + test2 + test4

			Df	Sum of Sq	RSS	AIC
<none>					554.10	85.462
- test4	1	59.74		613.84	86.021	
+ test3	1	10.79		543.31	86.970	
- test2	1	113.74		667.84	88.129	
- test1	1	411.30		965.40	97.342	